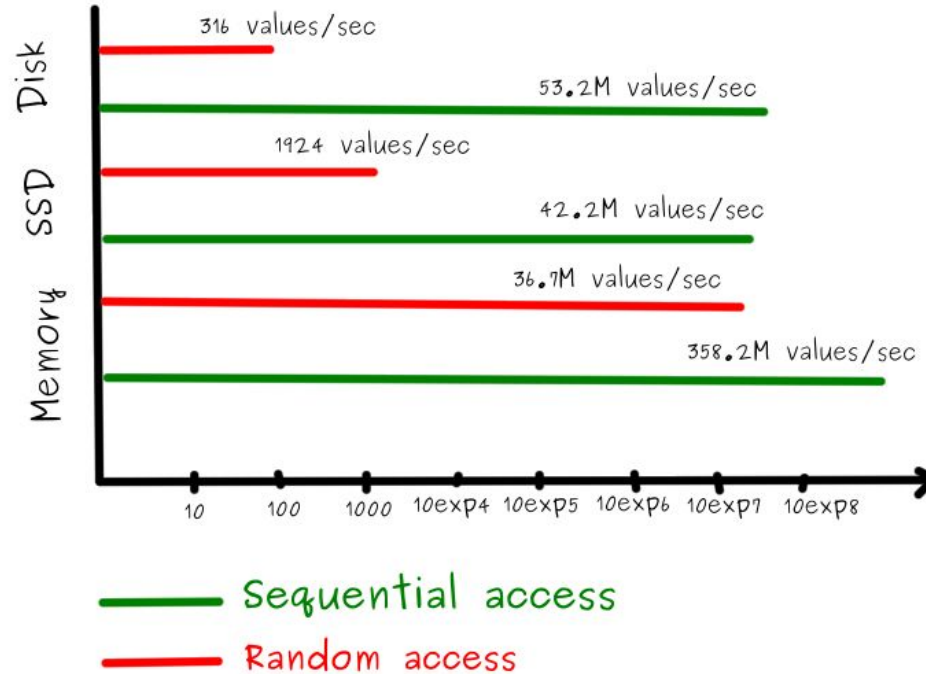


O-O-O-O-O-O-O

Persistence Axioms

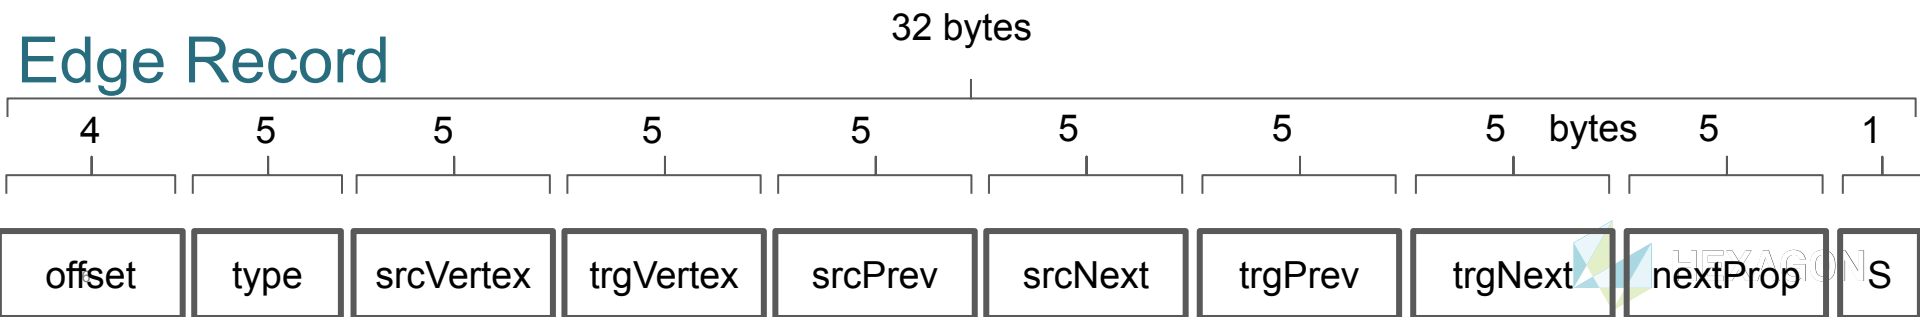


- *Sequential reads and writes are typically faster than random*
- *Random record access on collections of fixed-size records (eg. `get(index)`) can be very efficient* when based on computing access offsets using mathematical formulas (eg. $\text{offset} = \text{index} * \text{record-size}$) rather than scanning the data - ($O(1)$)
- Data retrieval optimization is typically a matter of *I/O optimization* (eg. minimize # of blocks loaded to resolve a given query)
- Data publishing optimization is typically a matter of *I/O optimization* (eg. minimize # of blocks to upload). Typically data publishing \neq data retrieval optimization.
- *Index-free adjacency is the key differentiator* of native graph processing (eg. edges to store direct pointers to adjacent nodes so that fast traversals are possible w/o the help of indices)

Sequential Access vs. Random Access

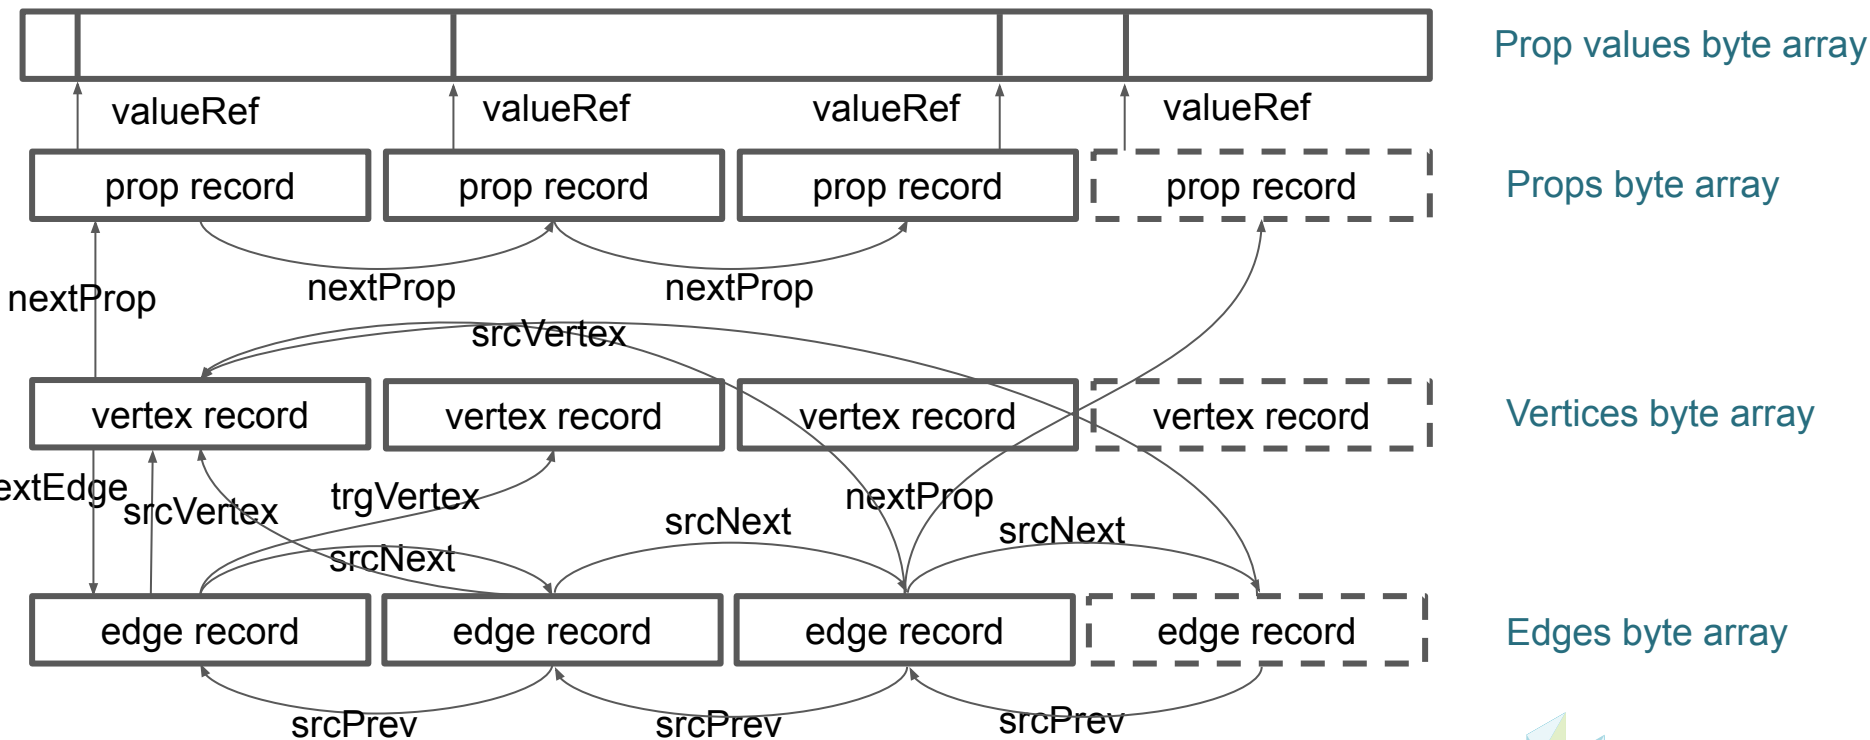


Benchmark published by Adam Jacobs, *The Pathologies of Big Data* in ACM Newsletter

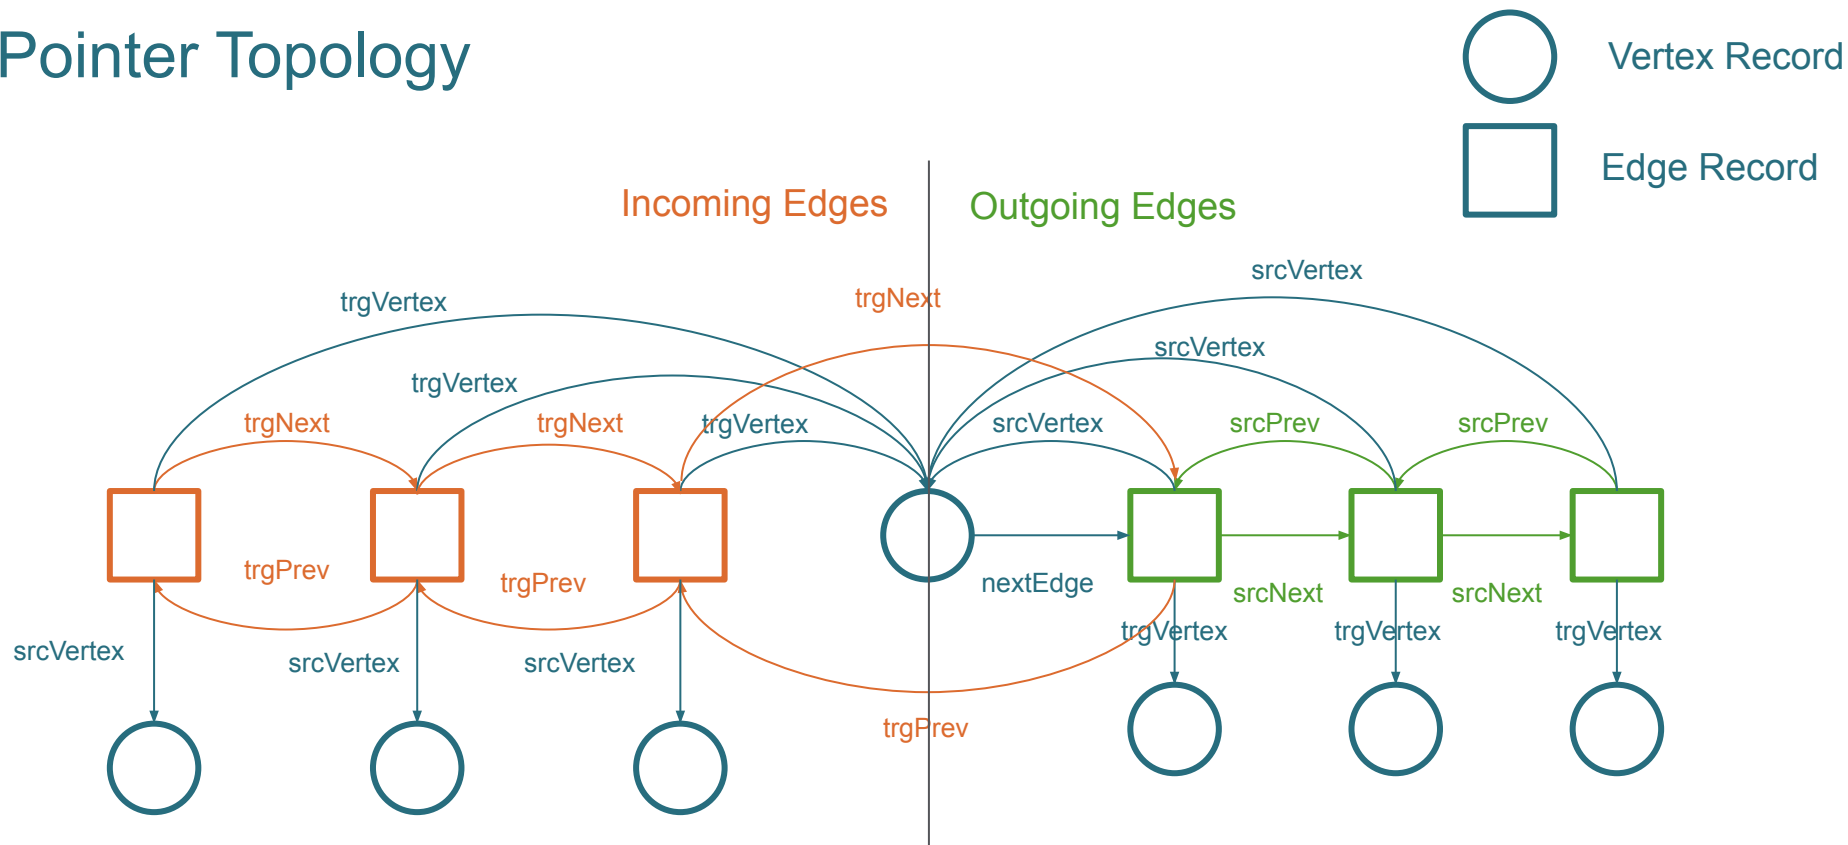
Persistence Format




Arrays of Records



Pointer Topology



Chunking



Stable Chunking: A method to partition any array of records to **support incrementality** (ie. generate stable chunks):

- Re-applying chunking alg. on the same data *should yield same chunks*
- Re-applying chunking alg. on a modified data set should *maximize reuse of older chunks and minimize the creation of new chunks*
- A modified chunk is considered new chunk (as it needs publishing)

Algorithm Choices:

- Fixed size chunking, on arrays of fixed size records (eg. chunk-size = $k \times$ record-size, where k = number of records per chunk)
- Content-based chunking, works on any data*

Index blocks by record offset



- The index describes unambiguously the array of records
- The index itself is an array of records which can be chunked recursively
- The index is immutable
- The index is content addressable

root cid	
0	block cid 0
offset 1	block cid 1
offset 2	block cid 2
offset 3	block cid 3
offset 4	block cid 4

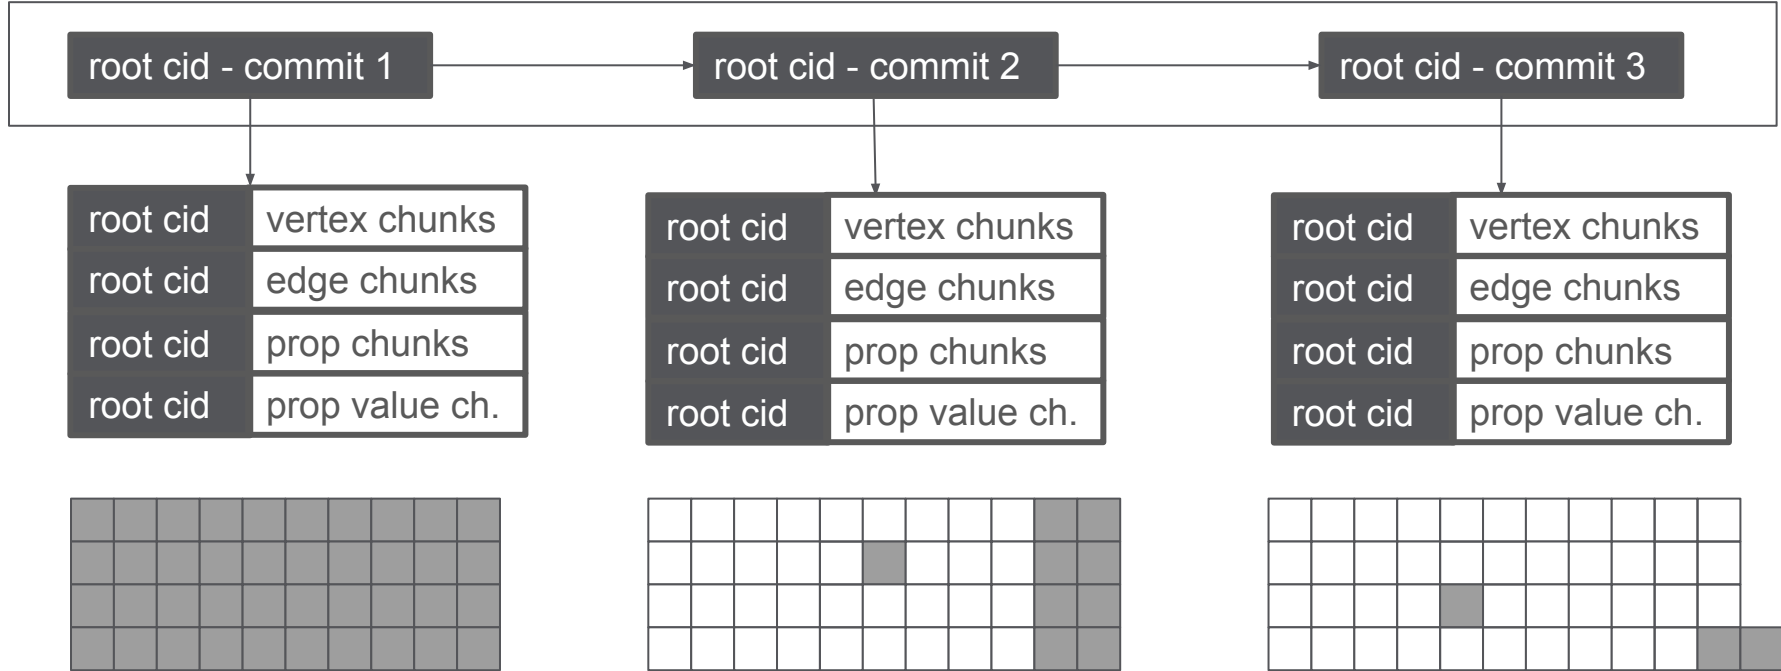
More Axioms

- *Content addressing* favors location independent storage, ie simplified replication across storage providers (ie. AZ blob storage, S3, browser local, etc.), supportive to collab. scenarios
- References represented as content identifiers *guarantee referred data integrity*: uniqueness, immutability and authenticity
- *Simplified & efficient block equality checks* w/o accessing referred data (eg. leverage the chunk index to identify changes to a given baseline)
- Supportive for chunk deduplication, supportive for versioning

Create Graph

Add new sub-tree

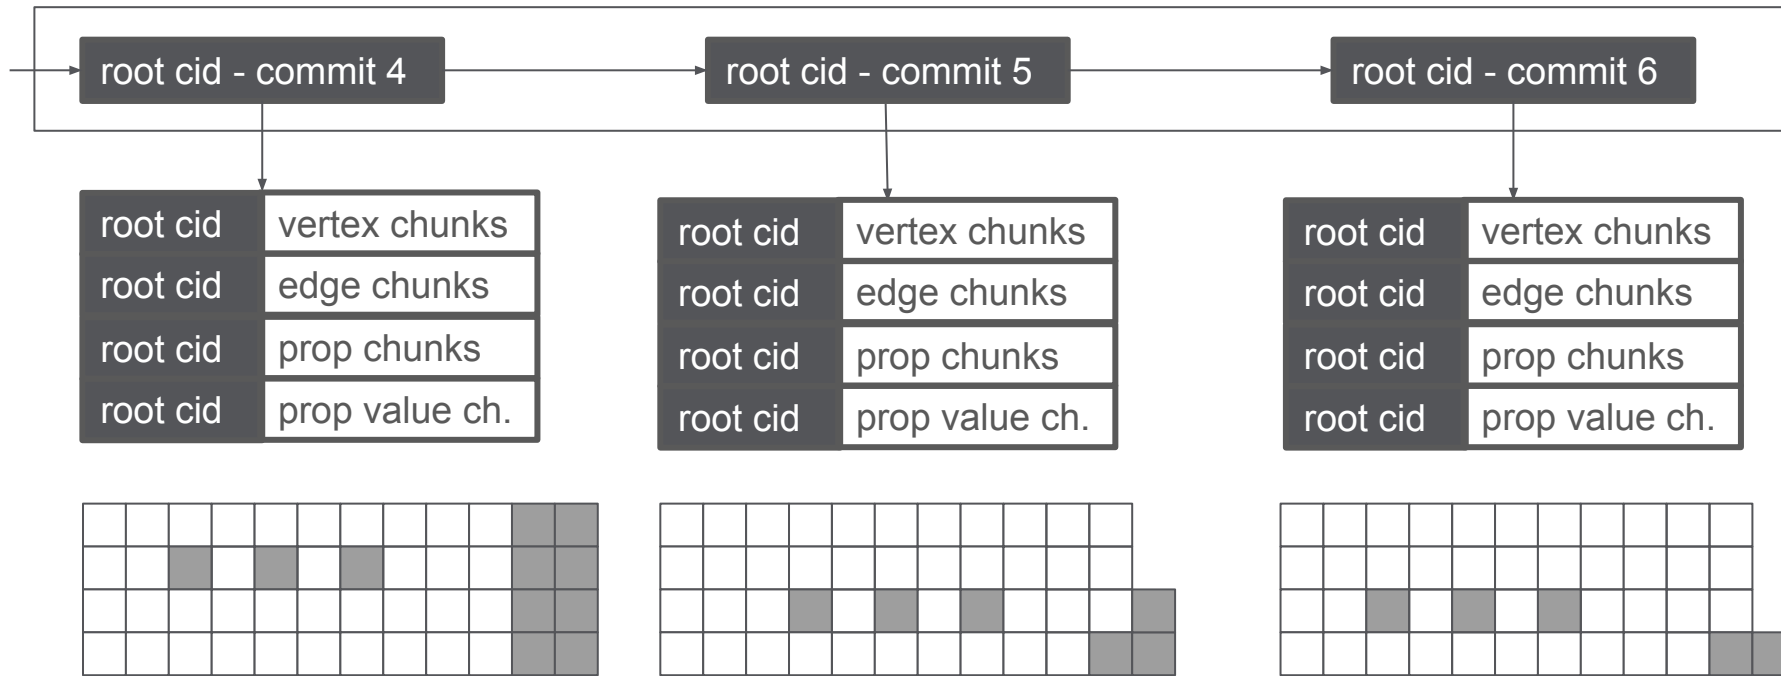
Modify many Property Values on single Vertex



Add many new sub-trees

Add many Properties
to many Vertices

Modify many Property
Values to many Vertices



Challenges



- **Incrementality vs. Fragmentation**

- Incrementality requires to collocate data according to edits and preserve historical blocks (minimize writes), but adds fragmentation to the actual graph structure (possibly scattered read). CQRS?

- **Read**

- Fragmentation can be a performance killer
- Use-case based optimizer (batching blocks at retrieval time)
- Caching blocks locally (relevant also for offline support)

- **Delete**

- Logical only, will benefit from a strategy to reuse offsets from deleted records

- **Compaction**

- To correct fragmentation, prune deleted records
- Invasive (potentially replace many offsets and associated refs, introduce logical offsets?),
- Complex to execute online